

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px

df = pd.read_csv("/content/winequality-red.csv")
df
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density
0	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.994
1	7.8	0.880	0.00	2.6	0.098	25.0	67.0	0.995
2	7.8	0.760	0.04	2.3	0.092	15.0	54.0	0.994
3	11.2	0.280	0.56	1.9	0.075	17.0	60.0	0.995
4	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.994
...
1594	6.2	0.600	0.08	2.0	0.090	32.0	44.0	0.995
1595	5.9	0.550	0.10	2.2	0.062	39.0	51.0	0.995
1596	6.3	0.510	0.13	2.3	0.076	29.0	40.0	0.995
1597	5.9	0.645	0.12	2.0	0.075	32.0	44.0	0.995
1598	6.0	0.310	0.47	3.6	0.067	18.0	42.0	0.995

1599 rows × 12 columns

Exploratory Data Analysis

```
print("rows,columns:"+str(df.shape))

rows,columns:(1599, 12)

df.head()
```

free


total

```
#missing variables  
print(df.isna().sum())
```

```
fixed acidity      0  
volatile acidity   0  
citric acid        0  
residual sugar     0  
chlorides          0  
free sulfur dioxide 0  
total sulfur dioxide 0  
density           0  
pH                0  
sulphates         0  
alcohol           0  
quality           0  
dtype: int64
```

Exploring data using variables


```
figure = px.histogram(df,x = 'quality')  
figure
```



```
#to find the correlation between variables
```

```
corr = df.corr()
plt.figure(figsize = (15,10))
sns.heatmap(corr, xticklabels = corr.columns,yticklabels = corr.columns,annot = True,cmap
```

<matplotlib.axes._subplots.AxesSubplot at 0x7eff69a7a990>



```
#classification version of target variable
df['goodquality'] = [1 if x>= 7 else 0 for x in df['quality']]
#seperate feature variables and target variable
X = df.drop(['quality','goodquality'],axis = 1)
y = df['goodquality']
```

```
#see proportion of good wines vs bad wines
df['goodquality'].value_counts()
```

```
0    1382
1     217
Name: goodquality, dtype: int64
```

```
#standardizing variables
from sklearn.preprocessing import StandardScaler
X_features = X
X = StandardScaler().fit_transform(X)
```

Split data

```
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size = 0.25,random_state = 0)

from sklearn.ensemble import RandomForestClassifier
classifier1 = RandomForestClassifier(n_estimators = 10, criterion = 'entropy', random_stat
classifier1.fit(X_train, y_train)
```

```
RandomForestClassifier(criterion='entropy', n_estimators=10, random_state=0)
```

```
from sklearn.metrics import confusion_matrix, accuracy_score
y_pred = classifier1.predict(X_test)
cm = confusion_matrix(y_test, y_pred)
print(cm)
accuracy_score(y_test, y_pred)
```

```
[[346   9]
 [ 23  22]]
0.92
```

```
from sklearn.svm import SVC
classifier2 = SVC(kernel = 'rbf', random_state = 0)
classifier2.fit(X_train, y_train)
```

```
SVC(random_state=0)
```

```
from sklearn.metrics import confusion_matrix, accuracy_score
y_pred = classifier2.predict(X_test)
cm = confusion_matrix(y_test, y_pred)
print(cm)
accuracy_score(y_test, y_pred)
```

```
[[349   6]
 [ 29  16]]
0.9125
```

```
import xgboost as xgb
model = xgb.XGBClassifier(random_state=1)
model.fit(X_train, y_train)
```


```
XGBClassifier(random_state=1)
```

```
from sklearn.metrics import confusion_matrix, accuracy_score
y_pred = model.predict(X_test)
cm = confusion_matrix(y_test, y_pred)
print(cm)
accuracy_score(y_test, y_pred)
```

```
[[336  19]
 [ 20  25]]
0.9025
```

```
feat_importances = pd.Series(classifier1.feature_importances_, index=X_features.columns)
feat_importances.nlargest(25).plot(kind='barh',figsize=(10,10))
```

<matplotlib.axes._subplots.AxesSubplot at 0x7eff58c6f990>



```
# Filtering df for only good quality
df_temp = df[df['goodquality']==1]
df_temp.describe()
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide
count	217.000000	217.000000	217.000000	217.000000	217.000000	217.000000	217.000000
mean	8.847005	0.405530	0.376498	2.708756	0.075912	13.981567	34.889400
std	1.999977	0.144963	0.194438	1.363026	0.028480	10.234615	32.572200
min	4.900000	0.120000	0.000000	1.200000	0.012000	3.000000	7.000000
25%	7.400000	0.300000	0.300000	2.000000	0.062000	6.000000	17.000000
50%	8.700000	0.370000	0.400000	2.300000	0.073000	11.000000	27.000000
75%	10.100000	0.490000	0.490000	2.700000	0.085000	18.000000	43.000000

```
# Filtering df for only bad quality
df_temp2 = df[df['goodquality']==0]
df_temp2.describe()
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide
count	1382.000000	1382.000000	1382.000000	1382.000000	1382.000000	1382.000000	1382.000000
mean	8.236831	0.547022	0.254407	2.512120	0.089281	16.172214	35.000000
std	1.682726	0.176337	0.189665	1.415778	0.049113	10.467685	32.572200
min	4.600000	0.160000	0.000000	0.900000	0.034000	1.000000	7.000000
25%	7.100000	0.420000	0.082500	1.900000	0.071000	8.000000	17.000000
50%	7.800000	0.540000	0.240000	2.200000	0.080000	14.000000	27.000000
75%	9.100000	0.650000	0.400000	2.600000	0.091000	22.000000	43.000000
max	15.900000	1.580000	1.000000	15.500000	0.611000	72.000000	138.000000



